

RESEARCH

Open Access



Evaluating data requirements for high-quality haplotype-resolved genomes for creating robust pangenome references

Prasad Sarashetti¹, Josipa Lipovac², Filip Tomas², Mile Šikić^{2,3*} and Jianjun Liu^{1,4*}

*Correspondence:
miles@gis.a-star.edu.sg;
liuj3@gis.a-star.edu.sg

¹ Laboratory of Human Genomics, Genome Institute of Singapore, A*STAR, Singapore, Singapore

² Laboratory for Bioinformatics and Computational Biology, Faculty of Electrical Engineering and Computing, University of Zagreb, Zagreb, Croatia

³ Laboratory of AI in Genomics, Genome Institute of Singapore, A*STAR, Singapore, Singapore

⁴ Department of Medicine, Yong Loo Lin School of Medicine, National University of Singapore, Singapore, Singapore

Abstract

Background: Long-read technologies from Pacific Biosciences (PacBio) and Oxford Nanopore Technologies (ONT) have transformed genomics research by providing diverse data types like HiFi, Duplex, and ultra-long ONT. Despite recent strides in achieving haplotype-phased gapless genome assemblies using long-read technologies, concerns persist regarding the representation of genetic diversity, prompting the development of pangenome references. However, pangenome studies face challenges related to data types, volumes, and cost considerations for each assembled genome, while striving to maintain sensitivity. The absence of comprehensive guidance on optimal data selection exacerbates these challenges.

Results: Our study evaluates recommended data types and volumes required to establish a robust de novo genome assembly pipeline for population-level pangenome projects, extensively examining performance between ONT's Duplex and PacBio HiFi datasets in the context of achieving high-quality phased genomes with enhanced contiguity and completeness. The results show that achieving chromosome-level haplotype-resolved assembly requires 20× high-quality long reads such as PacBio HiFi or ONT Duplex, combined with 15–20× of ultra-long ONT per haplotype and 10× of long-range data such as Omni-C or Hi-C. High-quality long reads from both platforms yield assemblies with comparable contiguity, with HiFi excelling in phasing accuracies, while Duplex generates more T2T contigs.

Conclusion: Our study provides insights into optimal data types and volumes for robust de novo genome assembly in population-level pangenome projects. Reassessing the recommended data types and volumes in this study and aligning them with practical economic limitations are vital to the pangenome research community, contributing to their efforts and pushing genomic studies with broader impacts.

Keywords: LRS special issue, Pangenome, De novo assembly, Sequencing platforms, Population-level studies



© The Author(s) 2024. **Open Access** This article is licensed under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License, which permits any non-commercial use, sharing, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons licence, and indicate if you modified the licensed material. You do not have permission under this licence to share adapted material derived from this article or parts of it. The images or other third party material in this article are included in the article's Creative Commons licence, unless indicated otherwise in a credit line to the material. If material is not included in the article's Creative Commons licence and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder. To view a copy of this licence, visit <http://creativecommons.org/licenses/by-nc-nd/4.0/>.

Background

A high-quality and complete human reference genome is the fundamental bedrock supporting genetic studies of human diseases and population structures. Over the past two decades, the human reference genome employed in genetic studies has been meticulously crafted from genomic segments sourced from thousands of individuals [1, 2]. Despite efforts to assemble high-quality, gapless genomes such as T2T-CHM13 [3], T2T-YAO [4], CN1 [5], I002C [6], or HG002 [7], such references raise concerns regarding their abilities to represent genetic variations across diverse human populations accurately. The prevailing consensus is that no singular reference sequence can adequately encapsulate the complex genomic diversity across global populations [8]. This understanding highlights the crucial need for high-quality reference genome panels that accurately resolve haplotypes, presenting the complex genetic variations observed within distinct populations [9–11]. In parallel, there is a growing trend to shift from a singular reference to a pangenomic approach, which supports a broader range of genomic diversity, acknowledging the complexities within and across diverse human populations [12–15]. This shift is supported by the rapid development of computational tools for pangenome construction and analysis [16–21].

Haplotype-resolved genome sequences are the building blocks for pangenome construction. However, despite the contradictory nature of cost and sensitivity, both of which play vital roles in pangenomic projects, most recent studies [14, 15] lack comprehensive evaluation and guidelines related to the optimal data types and volumes required, mostly relying on the propositions of assembly tool authors for volume and data type requirements.

At the forefront of long-read technology (LRT) innovation, Pacific Biosciences (PacBio) and Oxford Nanopore Technologies (ONT) stand out as the primary driving forces, spearheading advancements in this field through their groundbreaking contributions. PacBio's long reads (LR) have excelled in read quality, while ONT has leveraged its competitive edge in providing substantial read lengths at a lower cost [22]. To address the disparity between read quality and length, ONT has recently introduced a novel technique termed “Duplex,” capable of achieving a quality level of Q30, thereby bridging the gap between read quality and read length (<https://nanoporetech.com/about-us/news/oxford-nanopore-tech-update-new-duplex-method-q30-nanopore-single-molecule-reads-0>). Recent comparisons suggest that the two platforms exhibit similar performance in structural variation (SV) analysis [23, 24]. Details on the current LR sequencing platforms, LR mapping, variant calling, and genome assembly approaches are discussed elsewhere [25]. However, HiFi vs Duplex performance in genome assembly has not been properly evaluated and compared.

In this study, we evaluated different data types and the minimum data volume required to establish a robust pipeline of genome assembly for population-level pangenome projects. Specifically, we conducted a performance comparison between ONT's Duplex dataset and PacBio HiFi dataset. In this comparison, we extensively examined the performance of these datasets in the context of genome assembly, scrutinizing their effectiveness in achieving high-quality phased genomes with enhanced contiguity and completeness. Given the swift advancements in long-read technologies (LRT), it is prudent to reassess the recommended data types and volumes outlined in this study,

aligning them with the practical economic limitations within the scope of your research endeavors.

Results

DNA sequencing

The I002C data used for this research are generated as a part of an ongoing effort to generate telomere-2-telomere diploid assembly of a male Singaporean of Indian ancestry I002C [6]. Through sequencing on various platforms, we obtained the following dataset for the child sample: 152.97 Gb ($\sim 50.99\times$) PacBio HiFi data, 193.66 Gb ($\sim 64.55\times$) ONT Duplex data, 441.07 Gb ($\sim 147\times$) ONT Ultralong data (ULONT) and 222.21 Gb ($\sim 74.07\times$) Omni-C data. For the paternal sample, 107.69 Gb ($\sim 35.90\times$), and maternal sample, 112.48 Gb ($\sim 37.49\times$), MGI paired-end data was sequenced (Table 1). A similar volume of publicly available HG002 dataset was utilized in this study (Table 1). On average, the Duplex reads were twice as long as the HiFi reads, yet they maintained a comparable level of read quality (Fig. 1).

Coverage saturation analysis of population-scaled *de novo* assembly

To leverage the potential of long reads, for genome assembly we utilized high-quality long reads (HQLR) such as HiFi and Duplex, which are 10 kb or longer and ULONT reads of at least 100 kb (Additional file 2: Table S1). We examined the importance of diverse data types and offered general observations on the sequencing depth or data volume required for genome assembly and its analyses at scale.

Data down-sampling

We generated varying coverage depths by randomly down-sampling different data types, considering a haploid genome size of 3 Gbp (Additional file 2: Table S2–S4).

- i) HiFi and Duplex reads: downsampled datasets at $20\times$, $30\times$, $35\times$, $40\times$, and $45\times$ coverage
- ii) ULONT: downsampled datasets at $10\times$, $20\times$, $30\times$, $40\times$, $50\times$, and $60\times$ coverage
- iii) Omni-C/Hi-C: downsampled datasets at $10\times$, $20\times$, and $30\times$ coverage

Due to the longer length of Duplex reads, achieving the same sequencing depth requires, on average, twice as many HiFi reads as Duplex reads at any given coverage level, as demonstrated in Additional file 1: Fig. S1.

Evaluation of assembly results in terms of sequence saturation

To evaluate the impact of sequencing coverage on assembling performance and identify the coverage saturation point where assembly contiguity begins to plateau, we utilized hifiasm [26–28] to assemble HiFi/Duplex data independently (HQLR_Only) and in conjunction with ULONT data across varying coverage depths. The assembly results show a clear positive correlation between the augmentation of data coverage (HiFi/Duplex) and assembly performance for both primary assemblies (Additional file 1: Fig. S2), representing a mosaic of the two haplotypes and the two haplotypes derived from the phased assembly (Fig. 2).

Table 1 Summary of sequencing data from different platforms

Dataset	Sample	Platform	Data type	Reads	Total bases (bb)	Depth	Min Length	Max length	Average read length	N50	%GC
I002C	Child	PacBio	HiFi	8,986,857	152,977,790,522	50.99	8	63,921	17,022.40	17,132	40.78
		ONT	Duplex	6,561,950	193,663,321,924	64.55	1	191,644	29,513.10	38,128	40.85
		ONT	Ultralong	12,710,241	441,077,651,554	147.03	1	1,434,925	34,702.50	87,761	40.58
		MGI	Omni-C	1,476,892,046	222,216,283,752	74.07	150	150	150	150	42.5
I002A	Father	MGI	Short reads	717,979,504	107,696,925,600	35.9	150	150	150	150	40.48
I002B	Mother	MGI	Short reads	749,860,990	112,479,148,500	37.49	150	150	150	150	40.28
HG002	Child	PacBio	HiFi	9,076,876	164,744,547,100	54.91	86	63,894	18,149.92	17,963	40.32
		ONT	Duplex	6,110,824	147,724,801,341	49.24	1	187,925	24,174.29	34,205	40.89
		ONT	Ultralong	1,347,597	204,361,988,910	68.12	100,000	2,486,048	151,649	147,890	40.54
		Illumina	Hi-C	596,026,484	89,999,999,084	30	151	151	151	151	42.09
HG003	Father	Illumina	Short reads	687,439,454	101,741,039,192	33.91	148	148	148	148	39.67
HG004	Mother	Illumina	Short reads	684,524,092	101,309,565,616	33.77	148	148	148	148	39.95

The coverage depth calculation was based on a genome size estimation of 3 Gb. PacBio Pacific Biosciences, HiFi High-Fidelity reads, ONT Oxford Nanopore Technologies

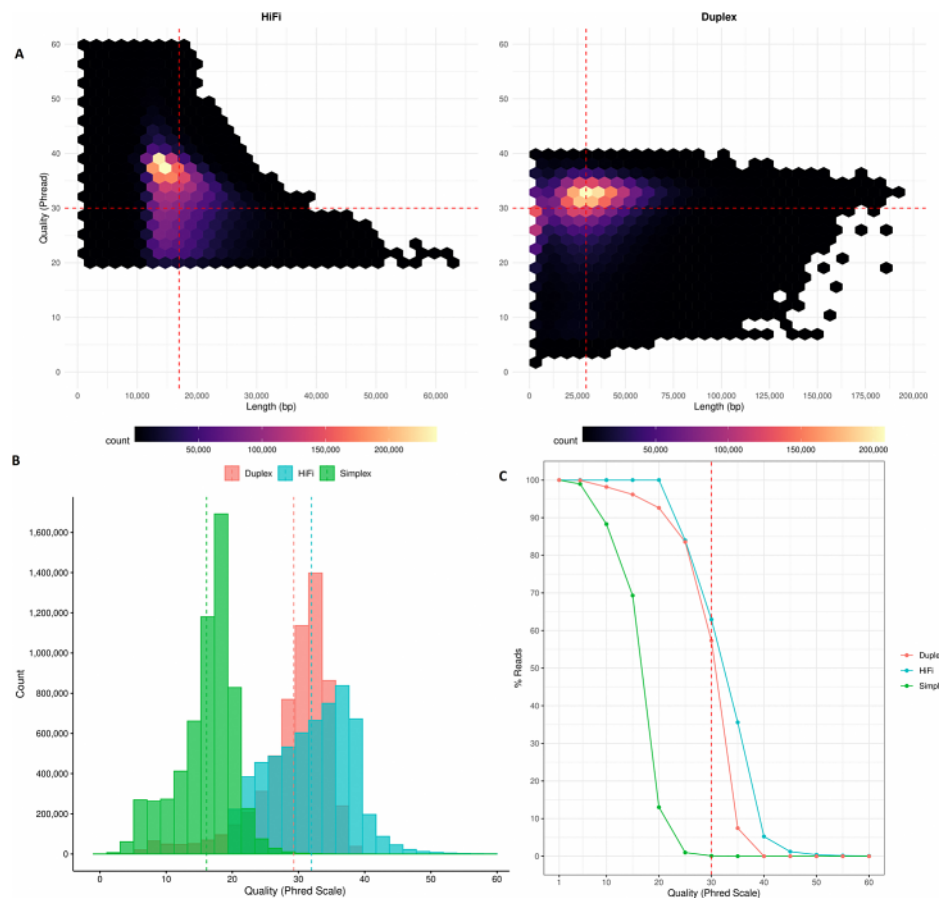


Fig. 1 Comparison of read length and quality (Phred scale) between PacBio HiFi and ONT Duplex reads. **A** Distribution of read length vs quality of HiFi and Duplex reads with vertical dotted lines indicating the average lengths: 17 kbp for HiFi and 29.5 kbp for Duplex reads. On average, more than 50% of both Duplex and HiFi reads have quality scores \geq Q30, a general cutoff for high-quality reads, indicated by the horizontal dotted line. **B** Comparison of read quality among ONT Simplex, ONT Duplex, and PacBio HiFi, with vertical dotted lines representing average quality scores: Q16 for Simplex, Q29 for Duplex, and Q32 for HiFi reads. **C** Percentage of reads with a quality score of Q30 and higher (dotted line). On average, 63% of HiFi reads and 57% of Duplex reads have a quality score of Q30 and higher

At any given coverage, the assembled genome size aligns well with expected genome sizes (2.9 Gb paternal, 3 Gb maternal, and 3.1 Gb primary assembly). The inflated assembled genome size positively correlated with the duplication rate (Rdup). As the data coverage increases, key assembly contiguity features such as NG50, Longest contig length, and Telomere-2-Telomere [T2T] contigs exhibit an upward trend, while the “No_of_Sequences” demonstrate a downward trajectory. Assembly contiguity reaches plateaus when the HQLR-only (HiFi/Duplex) coverage exceeds $35\times$ (Fig. 2).

Furthermore, in combination with ULONT data, even as low as $10\times$ ULONT along with $35\times$ of HQLR plateau coverage significantly enhances assembly contiguity compared to that of $45\times$ HQLR-only assembly. The inclusion of ULONT data notably improves the assembly of telomere-to-telomere contigs. Assembly contiguity reaches a plateau with ULONT coverage exceeding $30\times$. We observed a similar trend for primary assemblies (Additional file 1: Fig. S2). The detailed assembly statistics are provided

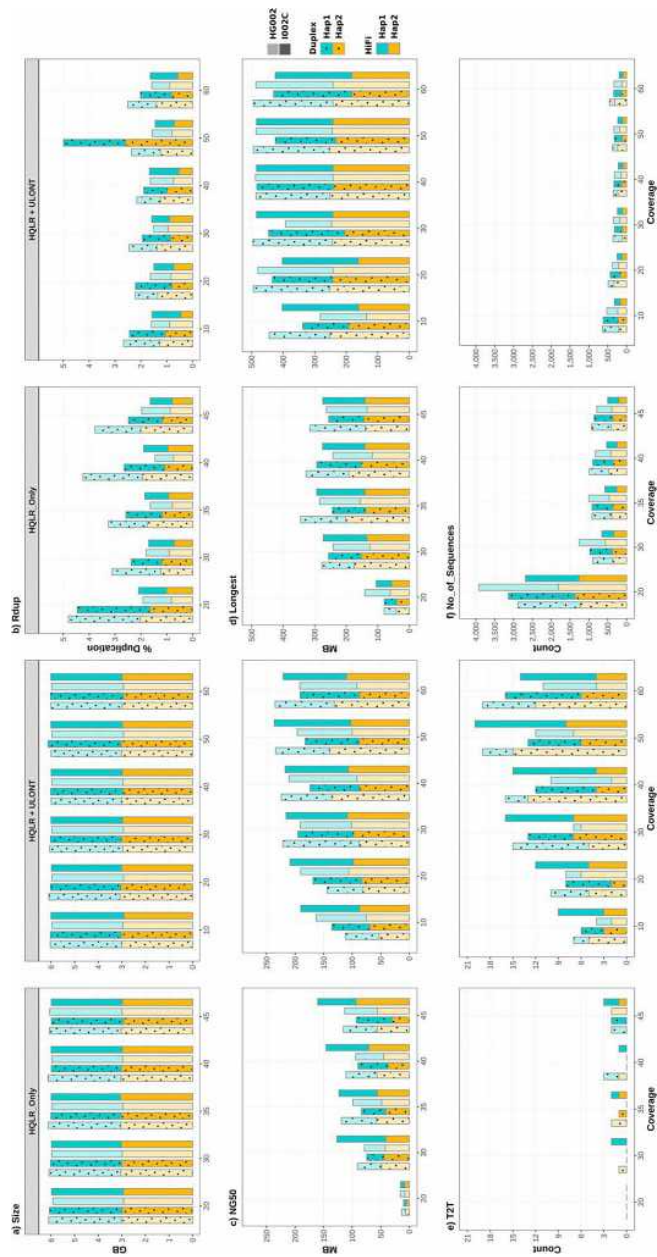


Fig. 2 Comparison of assembly performance versus data coverage. “HQLR Only” denotes assemblies generated solely with HiFi or Duplex data across various coverages, “HQLR + ULONT” signifies assemblies generated with a saturation coverage (35 x) of HiFi or Duplex data combined with various ULONT coverages

in the Additional Materials (Additional file 2: Table S5–S10). Hereafter, in the figures and tables, “HQLR + ULONT” denotes HQLR (HiFi/Duplex) coverage of $35\times$, representing the HQLR-only plateau coverage, combined with various ULONT coverages. Similarly, “HQLR + ULONT + Omni-C” signifies $35\times$ HQLR coverage combined with $30\times$ ULONT coverage, representing the ULONT plateau coverage, along with different levels of Omni-C (I002C)/Hi-C (HG002) coverage.

Improvement of phasing with Omni-C/Hi-C

The hifiasm tool is capable of producing pseudo-haplotypes or a dual assembly using HiFi/Duplex data alone (Fig. 3a) or in conjunction with ULONT (Fig. 3b). This process efficiently captures the heterozygous variances across the two haplotypes. HiFi-only assemblies demonstrate relatively fewer switch errors due to their higher quality compared to Duplex-only assemblies. Conversely, the longer read lengths of Duplex data contribute to achieving superior global phasing (hamming) compared to HiFi reads (Additional file 1: Fig. S3). Due to their lengths, ULONT reads additionally improve phasing [29]. However, even with ULONT reads, assemblers generate contigs with short phase blocks that often show increased phasing errors (Fig. 3b).

Incorporating even low coverage of long-range chromatin interaction data like Omni-C/Hi-C, such as $10\times$, results in a notable reduction in globally incorrectly phased variants (measured by hamming error), leveraging the long-range chromatin interaction information provided by Omni-C/Hi-C (Additional file 1: Fig. S3). Even though long-range interaction data can produce full-length phased contigs from different chromosomes, maternal and paternal origin contigs can be mixed in one haplotype (Fig. 3c). This intrinsic ambiguity in long-range interaction data phasing is attributed to the challenge of identifying markers that define paternal and maternal origin, a task not easily achievable with offspring data alone, except in the case of XY chromosomes. Despite this improvement, switch errors, which measure the local inaccuracies of heterozygous variants, remain largely unaffected due to the limitations in the information offered by long-range chromatin interaction data. Besides its phasing capability, Omni-C/Hi-C data can also be utilized for scaffolding. Omni-C coverage saturation concerning phasing can be observed at $10\times$ (Additional file 1: Fig. S4). Since hifiasm does not leverage long-range data for scaffolding to enhance contiguity (Additional file 1: Fig. S4), higher coverage may prove advantageous for scaffolding processes. Determining the optimal coverage for long-range chromatin interaction data (Omni-C/Hi-C) is beyond the scope of this study, as discussed elsewhere [30].

Genome completeness and quality

Genome completeness assessed through single-copy gene analysis revealed that assemblies from HQLR-only exhibited slightly lower performance with an average of 96.98% single copy, 1.40% duplicated, 0.30% fragmented, and 1.31% missing genes (Fig. 4a). Meanwhile, assemblies generated with HQLR + ULONT data showed higher completeness values with 97.53% single copy, 1.18% duplicated, 0.19% fragmented, and 1.10 missing genes (Fig. 4b). Combined haplotype results show increased coverage resulted in marginal improvements in gene completeness, with HQLR-only assemblies reaching

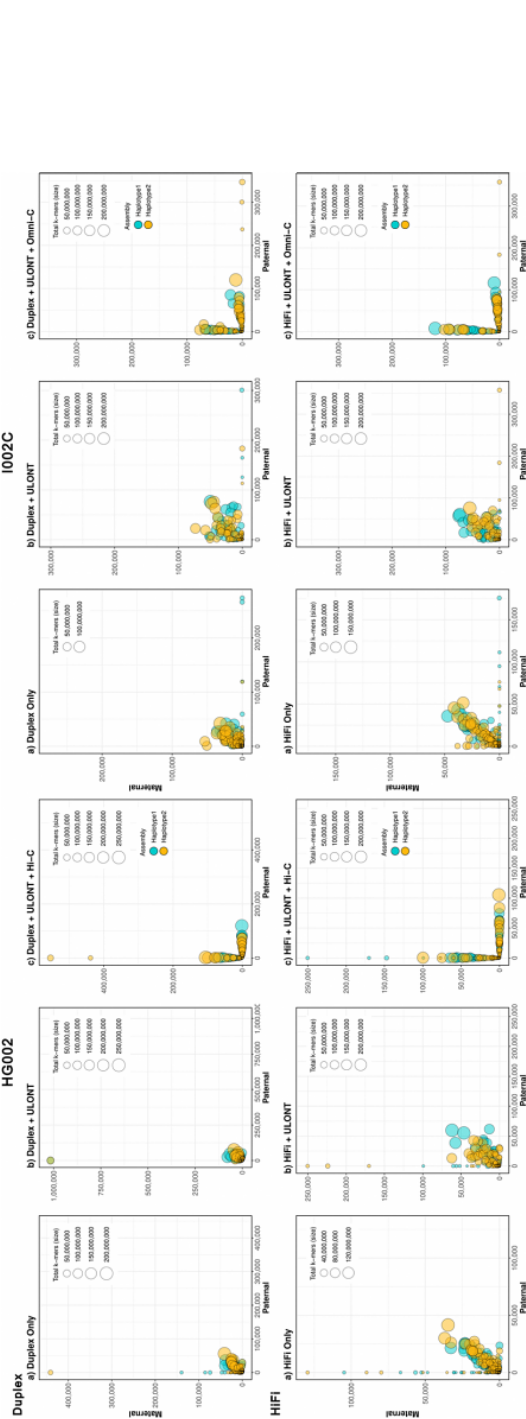


Fig. 3 Comparison of phasing accuracies of different assemblies (Duplex assemblies—top row, HIFI assemblies—bottom row). **a** Phasing accuracy of dual assembly generated from HQLR-only (HIFI/Duplex). **b** Phasing accuracy of dual assembly in conjunction with ULONT. **c** Haplotype separated assemblies with Omni-C/Hi-C data. Each circle denotes a contig, size reflecting its length. Circle's positions are determined by the number of maternal and paternal *k*-mers derived from high-quality short reads on respective contigs. Contigs positioned along the axis indicate higher phasing accuracy

saturation around $35 \times$ coverage, and ULONT assemblies around $30\text{--}40 \times$ (Fig. 4, Additional file 2: Table S11–12). Similar results were found for HG002 (Additional file 1: Fig. S5) and primary assemblies of both datasets (Additional file 1: Fig. S6–S7, Additional file 2: Table S13). Individual haplotypes from haplotype-resolved assemblies do not follow a clear trend of coverage saturation but show noticeable improvements by incorporating ULONT reads (Fig. 4, Additional file 1: Fig. S5).

The estimated k -mer completeness, which indicates the proportion of reliable k -mers from the reads found in the assembly, averaged 95.46% for haplotype-resolved assemblies (Fig. 5). In comparison, the primary assemblies averaged a slightly higher rate at 96.37 (Additional file 1: Fig. S8, Additional file 2: Table S14). This finding aligns with the gene completeness analysis results presented in (Fig. 4).

Assembly quality assessed from k -mers as measured by phred scale quality score (QV) generally showed improvement with increased coverage for both haplotype-resolved assemblies (Fig. 6, Additional file 2: Table S15) and primary assemblies (Additional file 1: Fig. S9, Additional file 2: Table S15).

Computational requirements

We conducted comparisons of both the runtime and peak memory consumption of assembling steps across various coverage levels for specific data types and assemblies resulting from the various combinations of different data types. The computational demands are of paramount importance, particularly in studies conducted at population scale and those utilizing cloud-based platforms for analysis.

The error correction process is a critical and most time-intensive step taking more than half of the total execution time, followed by the graph construction by long read assemblers. By default, hifiasm performs three rounds of error correction of input HiFi/Duplex reads. Consequently, the time and memory requirements exhibit an upward trajectory with increasing coverage when assemblies are derived solely from data (HQLR_only). In the case of “HQLR + ULONT,” where a fixed amount of HQLR data is employed, computational time shows an upward trend with the increased coverage of ULONT. At the same time, memory requirements remain stable across coverage levels. This stability in memory consumption is attributed to the implementation of ULONT data in their algorithm [28].

The incorporation of long-range chromatin interaction data (Omni-C) primarily utilized for phasing and resolving graph tangles reveals that both time and memory requirements remain more consistent across increased long-range data coverage (Fig. 7, Additional file 2: Table S16). However, compared to HQLR and “HQLR + ULONT,” the overall increase in memory requirements can be attributed to an additional step required to construct unique 31-mers from the initial assembly graph generated from HQLR reads for processing long-range data. This step depends on the coverage of HQLR, which remains fixed. In contrast, the variable coverage of long-range data has minimal influence on memory requirements, making it consistent across different levels of long-range data coverage. Computation time does not linearly increase with Hi-C coverage. Instead of using general-purpose read mappers to align Hi-C reads, hifiasm implements k -mer-based alignment to filter Hi-C reads that do not bridge heterozygous alleles or are






Fig. 5 Assessment of *k*-mer-based genome completeness analysis




Fig. 6 *K*-mer-based genome quality scores

mapped to homozygous unitigs reducing the computational burden as described in their algorithm [27].




Fig. 7 Computation resources consumed by hifiasm across different data types and coverages

Comparison of HiFi and Duplex reads performance inde novoassembly

We conducted a comparison between the new data types regarding their performance in genome assembly using the current I002C and publicly available HG002 datasets. HG002 Duplex and Revio HiFi data were downloaded from HPRC. The assemblies were constructed with the same coverage (i.e., $35 \times \text{HiFi/Duplex} + 30 \times \text{ULONT} + 10 \times \text{Hi-C [HG002]/Omni-C [I002C]}$) data with default parameters of hifiasm and Verkko [31] across three independent replicates. We assigned a rank of 1 to the highest value and 0 otherwise for each assembly feature. The sum of these ranks was then computed for both HiFi and Duplex assemblies to evaluate their performance based on specific criteria, ranging from best to worst (Fig. 8). Overall HiFi assemblies demonstrated lower values for metrics such as Rduplication, Number of Sequences, Switch, and Hamming errors, indicating superior assembly quality. However, Duplex assemblies achieved a higher count of T2T contigs and k -mer completeness. The lower NG50 of hifiasm Duplex assemblies may be due to hifiasm’s use of a string graph-based method, which struggles with handling contained reads, with read length being the aggravating factor. Since Duplex reads are, on average, twice as long, this issue is exacerbated. RAFT algorithm [32], minimizes this problem. However, the RAFT-hifiasm workflow requires RAFT to be executed once and hifiasm three times, making it at least two times slower than a single run of hifiasm. This issue is absent in de Bruijn graph-based assemblers like Verkko. The quantitative values for assembly features across replicates are available in the additional materials (Additional file 1: Fig. S10, Additional file 2: Table S17).

Discussion

The DNA sequencing landscape is continually evolving, with advancements in sequencing technologies offering unprecedented opportunities for genomic research. In this study, we conducted a comprehensive analysis of sequencing data obtained from PacBio HiFi, ONT Duplex, ONT Ultralong (ULONT), and Omni-C data in the context of genome assembly. We aimed to investigate coverage saturation for different data types and their implications for various aspects of genome assembly, including phasing, genome completeness, and assembly quality.

Our findings provide valuable insights into the optimal sequencing coverage depth required for genome assembly in large-scale analyses. Through coverage saturation analysis, we observed a positive correlation between sequencing coverage and assembly performance. Notably, assembly contiguity plateaued when the HQLR-only coverage exceeded $35\times$. Furthermore, the integration of ULONT data significantly enhanced assembly contiguity, particularly for assembling telomere-to-telomere contigs, underscoring the importance of long-range data in improving assembly contiguity. The assembly contiguity plateaus with ULONT coverage exceeding $30\times$.

We did not involve parental information in generating haplotype-resolved assemblies. Trio binning using parental data facilitates assembly and increases phasing accuracy compared to long-range data phasing [33]. However, it requires additional effort in the recruitment process and often parental information is not available. Even with high-quality long reads such as HiFi/Duplex and ULONT with substantial coverage, the assembled genome still can have higher switch and hamming errors. Our study demonstrates the efficacy of incorporating long-range chromatin interaction data like Omni-C/Hi-C to address this issue. By leveraging long-range contact information provided by long-range chromatin interaction data, we observed a notable reduction in globally incorrectly phased variants. However, challenges persist in accurately identifying the parental origin of phased contigs, highlighting the inherent ambiguity in long-range data phasing.

Genome completeness and quality assessments revealed marginal improvements with increased coverage, with assemblies incorporating ULONT data exhibiting higher completeness metrics compared to HQLR-only assemblies. Our analysis emphasizes the importance of considering both single-copy gene analysis and k -mer completeness for a comprehensive assessment of genome quality.

The computational demands associated with genomic analysis are substantial, particularly in population-scale studies. Our study highlights the time and memory requirements associated with the assembly process, emphasizing the need for efficient algorithms and computational resources to handle large datasets effectively.

As pioneers in long-read technology (LRT), PacBio and ONT continually refine their technologies and develop new advancements to deliver high-quality data at increasingly affordable prices. The recent launch of the PacBio Revio platform (<https://www.pacb.com/revio/>) stands as a testament to this commitment, elevating HiFi yield by $15\times$ while maintaining impeccable data quality compared to its predecessor, the PacBio Sequel II platform. The assembly contiguity achieved with HiFi data exhibits nearly identical performance on both the Sequel IIe and Revio platforms [24]. The substantial boost in data yield has effectively mitigated affordability concerns in comparison to competition.

Similarly, ONT has unveiled the enhanced R10 flowcell and introduced the innovative “Duplex” method, which achieves read quality nearing Q30 by sequencing both the template and complement strands of a single molecule. The effectiveness of these cutting-edge data types has been demonstrated in variant calling [24] and methylation studies [34], showcasing their utility and performance across different genomic applications.

A comparative analysis between PacBio HiFi and ONT Duplex data for genome assembly shows that HiFi data consistently delivers superior assembly quality, particularly in reducing duplication rates, sequence count, switch errors, and Hamming errors. However, Duplex data outperformed in producing a higher number of T2T contigs and *k*-mer completeness. Despite these strengths, both platforms exhibit comparable performance in terms of NG50 and the length of the longest contigs, highlighting that each method offers unique benefits depending on the specific assembly objective. These findings are backed by a recent study by Koren et al. [35], who evaluated ONT Duplex data for non-human samples and HG002.

Conclusion

Recognizing the dynamic nature of genomic research and the evolution of sequencing technologies and analytical methodologies is essential. Through our exploration of various sequencing data types and algorithms, we offer several key insights and recommendations for population-level pangenome reference generation efforts. We highlight the pivotal role of integrating high-quality data sources such as Pacbio HiFi/ONT Duplex and ONT ULONT, alongside long-range contract data like Omni-C, to achieve phased telomere-to-telomere level assemblies. In general, HiFi/Duplex coverage of $\geq 20\times$ complemented with $15\text{--}20\times$ of ULONT per haplotype and $10\times$ long-range data are essential requisites for attaining high-quality contiguous and phased assembly. We offer our findings as practical guidelines to help users choose sequencing platforms and coverage effectively.

Methods

Sample selection and preparation for sequencing

One family (comprising a mother, father, and child) with an Indian ethnic background out of 15 families recruited from Singapore as a part of the human genome project was selected. The selection criteria for the family were (1) current generation (child sample) is a male and (2) no genetic diseases with normal phenotype. All the participants were provided with informed consent for sample collection and usage including making data publicly available via databases. Sample collection and usage were approved by SingHealth Centralised Institutional Review Board. Whole blood was collected from the family (I002).

Isolation of peripheral blood mononuclear cells (PBMCs)

Ten milliliters of human whole blood samples were collected, and PBMCs were isolated using density gradient centrifugation with Ficoll-Paque (GE Healthcare). Blood was diluted with 20 ml of phosphate-buffered saline (PBS) and carefully layered over with 15 ml of Ficoll-Paque solution before centrifugation at 225 g for 30 min at room temperature. The PBMC layer was harvested, washed twice with PBS, and resuspended in




Fig. 8 The relative performance of HiFi vs Duplex assemblies from hifiasm and Verkko using I002C and HG002 dataset with plateau coverage, i.e., (35 × HiFi vs 35 × Duplex) + (30 × ULONT and 10 × Omni-C/Hi-C). For each assembly feature, we compared 24 assemblies obtained from 2 samples, 2 assemblies, 3 replicates, and 2 haplotypes, recording instances where the HiFi assembly outperformed the Duplex assembly and vice versa. “hifiasm + Verkko” shows the overall performance of HiFi vs Duplex across 24 assemblies. For example, 23 out of 24 HiFi assemblies had a lower switch error compared to Duplex assemblies, while 19 Duplex assemblies contained more T2T contigs than HiFi assemblies. Both platforms, however, produced comparable results in terms of NG50 and the length of the longest contigs

complete RPMI 1640 medium (Gibco) supplemented with 20% fetal bovine serum (FBS) and 1% penicillin–streptomycin.

Infection with B95-8 Epstein-Barr virus (EBV)

PBMCs were infected with the B95-8 strain of EBV by adding virus-containing supernatant derived from B95-8-infected marmoset B lymphocytes. The mixture was incubated at 37 °C with 5% CO₂ and left untouched for 8 days to facilitate virus entry into the B cells.

Cells were cryopreserved in a freezing medium containing 20% FBS and 10% dimethyl sulfoxide (DMSO), and stored in liquid nitrogen for long-term preservation.

Long-read sequencing (LRS) data generation

Pacbio data generation

The high molecular weight (HMW) DNA used for PacBio sequencing was extracted using the GentraPuregene kit (Qiagen; #158043) according to the manufacturer's instructions. Briefly, 1×10^7 frozen cell pellets from the I002C cell line were used as input for extraction. All vortexing steps were replaced with gentle inversion throughout the process, and 300 µl of Qiagen EB buffer was used for elution. Eluted DNA was incubated at 12 °C with gentle shaking over a period of 7 to 10 days. To avoid shearing the high molecular weight DNA, wide bore tips with gentle pipetting were used during handling. DNA was stored at 4 °C to prevent freeze and thaw cycle. Quantity and purity of extracted HMW DNA were assessed using triplicate concentration measurements from the top, middle, and bottom sections of sample volume, using Qubit dsDNA BR (Broad-Range) assay (ThermoFisher Scientific; Q32853) and NanoDrop 2000 spectrophotometer (ThermoFisher Scientific; ND-2000), according to manufacturer's instructions.

After DNA extraction, DNA fragment lengths were then measured using TapeStation 4200 (Agilent). Sequencing libraries were created using the SMRTbell Express Template Prep Kit 2.0 (PacBio) per the manufacturer's instructions. Libraries were sequenced on a Sequel IIe and Revio System (PacBio). After sequencing, CCS analyses were run using SMRTLink software v10 to produce HiFi reads.

ONT data generation

High molecular weight (HMW) gDNA extraction (Duplex sequencing) We obtained 12×10^6 frozen cell pellets from established lymphoblastoid cell line for the child sample and processed for HMW gDNA extraction using the Monarch HMW DNA Extraction Kit for Tissue (NEB; T3060). During the extraction, we excluded shaking during all incubation steps to preserve gDNA integrity. Quantity, purity, and integrity of extracted HMW gDNA were assessed using Qubit dsDNA BR assay (ThermoFisher Scientific; Q32853), NanoDrop 2000 spectrophotometer (ThermoFisher Scientific; ND-2000), and 15-h pulsed-field gel electrophoresis runs with the Pippin Pulse system (Sage Science; PPI0200), respectively. Quality-assessed HMW gDNA was then used for ligation-based

library preparation with the Ligation Sequencing Kit V14 (Oxford Nanopore Technologies; SQK-LSK114) to generate Duplex sequencing reads.

Ultra-high molecular weight (UHMW) gDNA extraction (ultra-long read sequencing) We processed 15×10^6 frozen cell pellets from an established lymphoblastoid cell line for the child sample for UHMW gDNA extraction using the Monarch HMW DNA Extraction Kit for Tissue, following the extraction steps described in the Ultra-Long DNA Sequencing Kit V14 (Oxford Nanopore Technologies; SQK-ULK114) protocol. Quality assessment of UHMW gDNA was performed similarly to HMW gDNA extraction. Quality-assessed UHMW gDNA was then used for transposase-based library preparation and purification with the Ultra-Long DNA Sequencing Kit V14 for the generation of ultra-long sequencing reads.

Library preparation and PromethION sequencing (Duplex and high duplex) We sheared 3 μg to 7.5 μg of extracted HMW gDNA to a target size of 55 kb to 60 kb and performed size-selective precipitation to remove DNA sizes < 25 kb. Repaired and end-prepped DNA was then used for library construction with SQK-LSK114 for both duplex and high duplex sequencing approaches. For standard duplex runs, libraries were loaded at 6 fmol to 7 fmol per load, while for high duplex runs, 7 fmol to 55 fmol of libraries were loaded and sequenced on PromethION 24 (Oxford Nanopore; PCA100024), R10.4.1 flowcells (Oxford Nanopore) FLO-PRO114M and FLO-PRO114HD respectively.

Library preparation and PromethION sequencing (ultra-long, UL) We used 40 μg to 45 μg of UHMW gDNA for ultra-long read library preparation using SQK-ULK114. Final UL libraries were sequenced on PromethION 24 using FLO-PRO114M flowcells with nuclease flushes performed at 23-h intervals.

The detailed steps of the entire procedure are outlined in the Additional Materials.

Omni-C data generation

The Dovetail Omni-C library was prepared using the Dovetail Omni-C™ Proximity Ligation Assay kit (Dovetail Genomics, Scotts Valley, CA, USA), according to the manufacturer's protocol (version 1.2). Briefly, after sample crosslinking with DSG and formaldehyde, chromatin was digested using a sequence-independent endonuclease and bound to chromatin capture beads. Proximity ligation was performed using a biotin-labeled bridge between the ends of the digested DNA. After reversal crosslinking, the DNA was purified and followed by library preparation. Finally, the biotinylated molecules were captured and amplified before sequencing on the Novaseq 6000 and HiSeq 4000 instruments (Illumina, San Diego, CA, USA) in paired-end mode.

Data analysis

All commands employed in the analysis are comprehensively listed in the Additional Materials file, providing readers with detailed procedures undertaken in this study.

Reads downsampling

To evaluate coverage saturation for both assembly contiguity and phasing efficiency, we downsampled the reads to various coverages. Reads were randomly subsampled to achieve the desired coverage utilizing Rasusa v0.7.1 [36], considering a genome size estimation of 3 gigabases (3 gb).

De novo assembly and assessment

The choice of assembler is critical for the evaluation process. We selected the assembler for coverage saturation analysis on the following criteria:

- 1) Ability to support different types of long reads
- 2) Native capability to generate haplotype-separated assemblies using a single data type and/or with additional data such as trio or long-range contact information
- 3) Computational demands
- 4) Active maintenance of the tool

Currently, the two most popular hybrid assemblers that support high-quality data, such as HiFi/Duplex, in addition to ultra-long (UL) reads, along with trio or long-range reads, to generate telomere-to-telomere haplotype-separated assemblies, are hifiasm [26–28] and Verkko [31]. However, when tested with the same dataset and computational configuration, Verkko's runtime was more than twice that of hifiasm (Additional file 1: Fig. S11, Additional file 2: Table S18). Furthermore, like other ONT assemblers, Verkko cannot produce haplotype-resolved assemblies using only Duplex or HiFi data unless the reads are first binned by haplotype for individual assembly or a diploid assembly is recovered from a haploid assembly using tools like HapDup. The recently published ONT assembler PECAT [37] can generate haplotype-wise assemblies from Duplex data alone. Still, it does not support the integration of additional datasets like ULONT or long-range interaction data. Given these limitations and insights from previous benchmark studies [38, 39], we selected hifiasm as the assembler to evaluate coverage saturation.

Assembly statistics

Assembly contiguity metrics were computed utilizing minigraph v0.20 [18] and paftools v2.26-r1175 [40].

Phasing statistics

The phasing efficiency of an assembly was evaluated in terms of switch error and Hamming error rates with Yak v0.1-r69-dirty [41] using parental short reads. Switch error quantifies the frequency of adjacent phased variants incorrectly transitioning between maternal and paternal haplotypes. Meanwhile, the Hamming error rate denotes the total misphased variants within each assembled contig. Phasing statistics were generated for both the haplotypes separately.

Assembly completeness and quality

To evaluate the impact of coverage variations on the completeness, we employed compleasm v0.2.2 [42] to obtain the BUSCO assessment results. Concurrently, we applied a k -mer-based approach for assembly completeness evaluation, using the KMC tool v3.2.1 [43]. Identifying reliable k -mers within the reads followed a previously outlined methodology [44]. The assembly completeness was computed as the fraction of reliable k -mers in the read set that also appeared in the assembly. Assembly QV was estimated using Yak.

Supplementary Information

The online version contains supplementary material available at <https://doi.org/10.1186/s13059-024-03452-y>.

Additional file 1. Supplemental methods and supplemental Figures S1–S11.

Additional file 2. Supplemental tables S1–S18.

Additional file 3. Review history.

Acknowledgements

The authors thank Oxford Nanopore for providing duplex and ULONT flow cells, as well as PacBio for HiFi flow cells at no cost. We would like to thank Mr Low Hwee Meng and Ms See Ting Leong from the Genome Institute of Singapore (GIS) Integrated Genomics Platform for their help with generating ONT data. Moreover, we would like to thank Mr Li Zhihui also from GIS Integrated Genomics Platform, for base calling the data as part of ONT-T2T collaboration. Finally, we would like to thank Ms Yao Fei for generating Omni-C libraries for the I002C T2T project and Ms Sia YY for maintaining the cell line, both from the GIS Laboratory of Human Genomics. This work was supported by the A*STAR Computational Resource Centre through the use of its high-performance computing facilities.

Peer review information

Editorial Board Member Shilpa Garg and Andrew Cosgrove were the primary editors of this article and managed its editorial process and peer review in collaboration with the rest of the editorial team.

Review history

The review history is available as Additional file 3.

Authors' contributions

J.J.L. and M.S. conceived the project. P.S. designed the pipeline. P.S. and J.J.L. prepared datasets and conducted the data analysis with the help of F.T. for genome completeness analysis. P.S. wrote the manuscript, and J.J.L. and F.T. helped with the organization of it. M.S. and J.J.L. supervised the project and provided mentorship.

Funding

This study was supported by the Agency for Science, Technology and Research (A*STAR) and the Singapore National Precision Medicine (NPM) Program (MOH-000588–01). Additional funding was provided by the European Union through the European Regional Development Fund under grant KK.01.1.1.01.0009 (DATACROSS), as well as by the Croatian Science Foundation under grant IP-2018–01-5886 (SIGMA).

Data availability

The I002C data utilized in this research are generated as part of an ongoing initiative to develop a telomere2-telomere diploid assembly of I002C [6]. The reads are submitted to NCBI under project ID PRJNA1150503. Yak files for parental data, used for assembly phasing analysis are added to zenodo (<https://doi.org/https://doi.org/10.5281/zenodo.14242314> [45]). Additionally, the HG002 dataset used in this study is accessible via AWS from the Human Pangenome Reference Consortium (HPRC).

Pacbio Revio HiFi data downloaded from: https://human-pangenomics.s3.amazonaws.com/index.html?prefix=submissions/80d00e88-7a92-46d8-88c7-48f1486e11ed--HG002_PACBIO_REVIO/.

ONT Duplex data downloaded from: https://human-pangenomics.s3.amazonaws.com/index.html?prefix=submissions/0CB931D5-AE0C-4187-8BD8-B3A9C9BFDADE--UCSC_HG002_R1041_Duplex_Dorado/Dorado_v0.1.1/.

ONT ultra long data downloaded from: https://s3-us-west-2.amazonaws.com/human-pangenomics/index.html?prefix=NHGRI_UCSC_panel/HG002/nanopore/ultra-long/.

Hi-C data downloaded from: https://s3-us-west-2.amazonaws.com/human-pangenomics/index.html?prefix=NHGRI_UCSC_panel/HG002/hic/.

Declarations

Ethics approval and consent to participate

All the participants were provided with informed consent for sample collection, and usage including making data publicly available via databases. Sample collection and usage were approved by SingHealth Centralised Institutional Review Board (IRB Reference: 2024–069). All experimental methods comply with the Helsinki Declaration.

Consent for publication

Not applicable.

Competing interests

M.S. has been jointly funded by Oxford Nanopore Technologies and AI Singapore for the project AI-driven De Novo Diploid Assembler. The remaining authors declare no competing interests.

Received: 29 March 2024 Accepted: 29 November 2024

Published online: 18 December 2024

References

1. International Human Genome Sequencing Consortium, Lander ES, Linton LM, Birren B, Nusbaum C, Zody MC, et al. Initial sequencing and analysis of the human genome. *Nature*. 2001;409:860–921. <https://doi.org/10.1038/35057062>.
2. Schneider VA, Graves-Lindsay T, Howe K, Bouk N, Chen H-C, Kitts PA, et al. Evaluation of GRCh38 and de novo haploid genome assemblies demonstrates the enduring quality of the reference assembly. *Genome Res*. 2017;27:849–64. <https://doi.org/10.1101/gr.213611.116>.
3. Nurk S, Koren S, Rhie A, Rautiainen M, Bizikadze AV, Mikheenko A, et al. The complete sequence of a human genome. *Science*. 2022;376:44–53. <https://doi.org/10.1126/science.abj6987>.
4. He Y, Chu Y, Guo S, Hu J, Li R, Zheng Y, et al. T2T-YAO: a telomere-to-telomere assembled diploid reference genome for Han Chinese. *Genomics Proteomics Bioinformatics*. 2023. <https://doi.org/10.1016/j.gpb.2023.08.001>.
5. Yang C, Zhou Y, Song Y, Wu D, Zeng Y, Nie L, et al. The complete and fully-phased diploid genome of a male Han Chinese. *Cell Res*. 2023;33:745–61. <https://doi.org/10.1038/s41422-023-00849-5>.
6. Lab of Human Genomics. I002C: Telomere-to-Telomere diploid Indian Genome. Available from: <https://github.com/LHG-GG/I002C> or <https://github.com/lbcb-sci/I002C>. Accessed 24 Mar 2024.
7. Telomere-to-Telomere (T2T) consortium. HG002: A complete diploid human genome. Available from: <https://github.com/marbl/HG002>. Accessed 24 Mar 2024.
8. Yang X, Lee W-P, Ye K, Lee C. One reference genome is not enough. *Genome Biol*. 2019;20:104. <https://doi.org/10.1186/s13059-019-1717-0>.
9. Lou H, Gao Y, Xie B, Wang Y, Zhang H, Shi M, et al. Haplotype-resolved de novo assembly of a Tujia genome suggests the necessity for high-quality population-specific genome references. *Cell Syst*. 2022;13:321–333.e6. <https://doi.org/10.1016/j.cels.2022.01.006>.
10. Deng L, Xie B, Wang Y, Zhang X, Xu S. A protocol for applying a population-specific reference genome assembly to population genetics and medical studies. *STAR Protoc*. 2022;3:101440. <https://doi.org/10.1016/j.xpro.2022.101440>.
11. Ballouz S, Dobin A, Gillis JA. Is it time to change the reference genome? *Genome Biol*. 2019;20:20. <https://doi.org/10.1186/s13059-019-1774-4>.
12. Sherman RM, Salzberg SL. Pan-genomics in the human genome era. *Nat Rev Genet*. 2020;21:243–54. <https://doi.org/10.1038/s41576-020-0210-7>.
13. Wang T, Antonacci-Fulton L, Howe K, Lawson HA, Lucas JK, Phillippy AM, et al. The human pangenome project: a global resource to map genomic diversity. *Nature*. 2022;604:437–46. <https://doi.org/10.1038/s41586-022-04601-8>.
14. Liao W-W, Asri M, Ebler J, Doerr D, Haukness M, Hickey G, et al. A draft human pangenome reference. *Nature*. 2023;617:312–24. <https://doi.org/10.1038/s41586-023-05896-x>.
15. Gao Y, Yang X, Chen H, Tan X, Yang Z, Deng L, et al. A pangenome reference of 36 Chinese populations. *Nature*. 2023;619:112–21. <https://doi.org/10.1038/s41586-023-06173-7>.
16. Garrison E, Sirén J, Novak AM, Hickey G, Eizenga JM, Dawson ET, et al. Variation graph toolkit improves read mapping by representing genetic variation in the reference. *Nat Biotechnol*. 2018;36:875–9. <https://doi.org/10.1038/nbt.4227>.
17. Eizenga JM, Novak AM, Sibbesen JA, Heumos S, Ghaffari A, Hickey G, et al. Pangenome graphs. *Annu Rev Genomics Hum Genet*. 2020;21:139–62. <https://doi.org/10.1146/annurev-genom-120219-080406>.
18. Li H, Feng X, Chu C. The design and construction of reference pangenome graphs with minigraph. *Genome Biol*. 2020;21:21. <https://doi.org/10.1186/s13059-020-02168-z>.
19. Vernikos GS. A review of pangenome tools and recent studies. *The Pangenome*. Cham: Springer International Publishing; 2020. p. 89–112.
20. Ebler J, Ebert P, Clarke WE, Rausch T, Audano PA, Houwaart T, et al. Pangenome-based genome inference allows efficient and accurate genotyping across a wide spectrum of variant classes. *Nat Genet*. 2022;54:518–25. <https://doi.org/10.1038/s41588-022-01043-w>.
21. Hickey G, Monlong J, Ebler J, Novak AM, Eizenga JM, Gao Y, et al. Pangenome graph construction from genome alignments with minigraph-cactus. *Nat Biotechnol*. 2023;42:663. <https://doi.org/10.1038/s41587-023-01793-w>.
22. Logsdon GA, Vollger MR, Eichler EE. Long-read human genome sequencing and its applications. *Nat Rev Genet*. 2020;21:597–614. <https://doi.org/10.1038/s41576-020-0236-x>.
23. Du X, Li L, Liang F, Liu S, Zhang W, Sun S, et al. Robust benchmark structural variant calls of an Asian using state-of-the-art long-read sequencing technologies. *Genomics Proteomics Bioinformatics*. 2022;20:192–204. <https://doi.org/10.1016/j.gpb.2020.10.006>.
24. Harvey WT, Ebert P, Ebler J, Audano PA, Munson KM, Hoekzema K, et al. Whole-genome long-read sequencing downsampling and its effect on variant calling precision and recall. *bioRxiv*. 2023. <https://doi.org/10.1101/2023.05.04.539448>.
25. De Coster W, Weissensteiner MH, Sedlazeck FJ. Towards population-scale long-read sequencing. *Nat Rev Genet*. 2021;22:572–87. <https://doi.org/10.1038/s41576-021-00367-3>.

26. Cheng H, Concepcion GT, Feng X, Zhang H, Li H. Haplotype-resolved de novo assembly using phased assembly graphs with hifiasm. *Nat Methods*. 2021;18:170–5. <https://doi.org/10.1038/s41592-020-01056-5>.
27. Cheng H, Jarvis ED, Fedrigo O, Koepfli K-P, Urban L, Gemmell NJ, et al. Haplotype-resolved assembly of diploid genomes without parental data. *Nat Biotechnol*. 2022;40:1332–5. <https://doi.org/10.1038/s41587-022-01261-x>.
28. Cheng H, Asri M, Lucas J, Koren S, Li H. Scalable telomere-to-telomere assembly for diploid and polyploid genomes with double graph. *Nat Methods*. 2024;21:967–70. <https://doi.org/10.1038/s41592-024-02269-8>.
29. Jain M, Koren S, Miga KH, Quick J, Rand AC, Sasani TA, et al. Nanopore sequencing and assembly of a human genome with ultra-long reads. *Nat Biotechnol*. 2018;36:338–45. <https://doi.org/10.1038/nbt.4060>.
30. Sur A, Noble WS, Myler PJ. A benchmark of Hi-C scaffolders using reference genomes and *de novo* assemblies. *bioRxiv*. 2022. <https://doi.org/10.1101/2022.04.20.488415>.
31. Rautiainen M, Nurk S, Walenz BP, Logsdon GA, Porubsky D, Rhie A, et al. Telomere-to-telomere assembly of diploid chromosomes with Verkko. *Nat Biotechnol*. 2023;41:1474–82. <https://doi.org/10.1038/s41587-023-01662-6>.
32. Kamath SS, Bindra M, Pal D, Jain C. Telomere-to-telomere assembly by preserving contained reads. *bioRxiv*. 2023. <https://doi.org/10.1101/2023.11.07.565066>.
33. Koren S, Rhie A, Walenz BP, Dilthey AT, Bickhart DM, Kingan SB, et al. De novo assembly of haplotype-resolved genomes with trio binning. *Nat Biotechnol*. 2018;36:1174–82. <https://doi.org/10.1038/nbt.4277>.
34. Ni Y, Liu X, Simeneh ZM, Yang M, Li R. Benchmarking of nanopore R10.4 and R9.4.1 flow cells in single-cell whole-genome amplification and whole-genome shotgun sequencing. *Comput Struct Biotechnol J*. 2023;21:2352–64. <https://doi.org/10.1016/j.csbj.2023.03.038>.
35. Koren S, Bao Z, Guarracino A, Ou S, Goodwin S, Jenike KM, et al. Gapless assembly of complete human and plant chromosomes using only nanopore sequencing. *bioRxiv*. 2024. <https://doi.org/10.1101/2024.03.15.585294>.
36. Hall M. Rasusa: Randomly subsample sequencing reads to a specified coverage. *J Open Source Softw*. 2022;7:3941. <https://doi.org/10.21105/joss.03941>.
37. Nie F, Ni P, Huang N, Zhang J, Wang Z, Xiao C, et al. De novo diploid genome assembly using long noisy reads. *Nat Commun*. 2024;15:1–15. <https://doi.org/10.1101/2022.09.25.509436>.
38. Jarvis ED, Formenti G, Rhie A, Guarracino A, Yang C, Wood J, et al. Semi-automated assembly of high-quality diploid human reference genomes. *Nature*. 2022;611:519–31. <https://doi.org/10.1038/s41586-022-05325-5>.
39. Li H, Durbin R. Genome assembly in the telomere-to-telomere era. *arXiv [q-bio.GN]*. 2023. <https://doi.org/10.48550/arXiv.2308.07877>.
40. Li H. Minimap2: pairwise alignment for nucleotide sequences. *Bioinformatics*. 2018;34:3094–100. <https://doi.org/10.1093/bioinformatics/bty191>.
41. Li H. yak: Yet another k-mer analyzer. Available from: <https://github.com/lh3/yak>. Accessed 24 Mar 2024.
42. Huang N, Li H. compleasm: a faster and more accurate reimplement of BUSCO. *Bioinformatics*. 2023;39:39. <https://doi.org/10.1093/bioinformatics/btad595>.
43. Kokot M, Długosz M, Deorowicz S. KMC 3: counting and manipulating *k*-mer statistics. *Bioinformatics*. 2017;33:2759–61. <https://doi.org/10.1093/bioinformatics/btx304>.
44. Rhie A, Walenz BP, Koren S, Phillippy AM. Merqury: reference-free quality, completeness, and phasing assessment for genome assemblies. *Genome Biol*. 2020;21:21. <https://doi.org/10.1186/s13059-020-02134-9>.
45. Sarashetti P. parental yak files. <https://doi.org/10.5281/zenodo.14242314>.

Publisher's Note

Springer Nature remains neutral with regard to jurisdictional claims in published maps and institutional affiliations.